

BOĞAZİÇİ UNIVERSITY

CMPE 300

ANALYSIS OF ALGORITHMS

**MapReduce Algorithm
MPI Programming Project**

Mustafa Enes ÇAKIR

2013400105

December 17, 2017

Submitted to Mehmet Köse

Contents

1	Introduction	2
2	Program Interface	2
3	Program Execution	3
4	Input and Output	3
5	Program Structure	4
5.1	Reading and Dividing	4
5.2	Mapping	4
5.3	Sorting	4
5.4	Merging	4
5.5	Reducing	5
5.6	Word Struct	5
6	Examples	6
7	Improvements and Extensions	6
8	Difficulties Encountered	6
9	Conclusion	6
10	Appendices	7

1 Introduction

We are asked to implement MapReduce algorithm using Message Passing Interface (MPI) for calculating word frequencies. Master processor coordinates other slave processors. After this point I will use *master* for master processor and *slaves* for slave processors. *Master* shares data between *slaves* in balanced parts and combines results.

In the given MapReduce algorithm, there 5 basic steps. First of all, *master* reads input file and send raw string lists to *slaves*. Secondly, *slaves* map words with their counts. In this step, we don't look other words, just consider current one. Evidently each word is mapped with count 1. *Slaves* sends mapped word lists to *master* and it combines these list. Thirdly *master* divides mapped word list to *slaves*. They sort corresponding lists lexically and send them back to *master*. Fourthly, *Master* merges sorted list. Finally, it reduces word with summing up their counts.

2 Program Interface

It's a basic command line application written with C++ using Boost MPI library. For compiling and running application, Boost MPI and Open MPI have to be installed. Boost libraries have to be added to *mpic++* that is Open MPI C++ compiler. *Figure 1* shows compile command in detail.

```
mpic++ word_freq.cpp -o word_freq.o -std=c++11 -I/usr/local/include/ -L/usr/local/lib -lboost_mpi -lboost_serialization
```

Source Code Output Program C++ Version Boost MPI Libraries

Figure 1: Compilation Command

Compiled code can be executed with *mpirun* command with given number of processor. Application takes 2 parameters. First one is input file, and second one is output file. If output file isn't provides, result printed to console.

```
mpirun -np 4 word_freq.o input.txt output.txt
```

Number of processor Compiled Code Input File Output File

Figure 2: Run Command

3 Program Execution

The command line application takes 2 arguments. First one is input file and it's required. Second one is output file and it's optional. If no file is provided, application prints out output to console. It calculates frequencies of words in given input file, and print words with their frequency in dictionary order. You can see sample execution of application in *Figure 3*.



```
→ CmpE300 cat test.txt
sample
test
sample
august
august
august
computer
sample
engineer
car

→ CmpE300 mpirun -np 4 word_freq.o test.txt
august 3
car 1
computer 1
engineer 1
sample 3
test 1
```

Figure 3: Program Execution with sample input

4 Input and Output

The input has one word on each line. In the other word, it's in tokenized format and has no limitations.

```
sample
test
sample
august
august
august
computer
sample
engineer
car
```

Listing 1: Sample input

Output has one word and it's count on each line in lexical order. It might be in file or on console. If output file name is provided, it prints out to file.

```
august 3
car 1
computer 1
engineer 1
sample 3
test 1
```

Listing 2: Sample output of the sample input

5 Program Structure

The application is developed with C++ with Boost MPI. I used *Word* struct for data representation. I choose rank 0 processor as *master*.

5.1 Reading and Dividing

Master checks arguments of application. If no input file is given, it raises an error. If an output file is given, changes standard output with file. It reads input file line by line and pushes word strings to a vector. *Master* divides this vector using index modulo. For example words indexed by 2, 5, 8, 11 goes to *slave* rank with 2 in world that is sized 4.

5.2 Mapping

Slaves get their string subvectors and constructs *Word* structs with count 1 for every word string. Lastly *slaves* send mapped words to *master*. *Master* combines coming *Word* vectors.

5.3 Sorting

Master shares this mapped vector to *slaves* again. *Slaves* sort received subvectors using *std::sort* and send back to *master*.

5.4 Merging

So received vectors by master are sorted now. As the last step of merge sort, *master* merges that sorted vectors with *std::merge*. For this merge opera-

tion I used *temp_words* vector. After merging to *temp_words*, I swapped *temp_words* with *sorted_words*.

5.5 Reducing

Master traverses *sorted_words* vector and reduces it. If the current *Word* is same as the last word, increase count of the last word. Else, it changes the last word with current one.

5.6 Word Struct

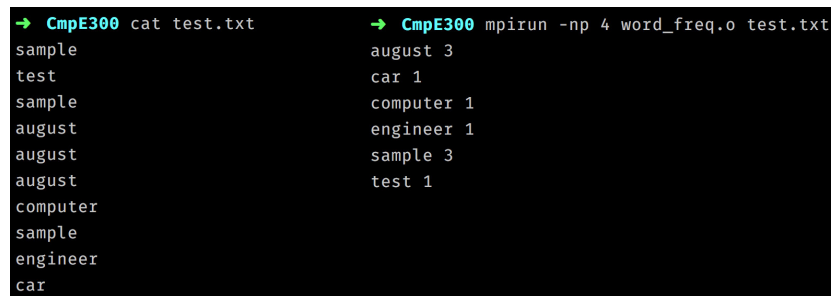
The main data structure in application is *Word* struct. It's combination of word's text and it's count. Unusual part of this code is *serialize* method. It's transform the word to format that can be send via MPI.

```
1 struct Word {
2 public:
3     string text;
4     int count;
5
6 private:
7     friend class boost::serialization::access;
8
9     /**
10      * Serializes Word struct for sending via Boost MPI
11      */
12     template<class Archive>
13     void serialize(Archive &ar, const unsigned int version) {
14         ar & text;
15         ar & count;
16     }
17
18 };
```

Code 1: Code of Word struct

6 Examples

The application executed with sample input *Listing 1*. It gives sample output *Listing 2*.



```
→ CmpE300 cat test.txt          → CmpE300 mpirun -np 4 word_freq.o test.txt
sample                          august 3
test                            car 1
sample                          computer 1
august                          engineer 1
august                          sample 3
august                          test 1
computer
sample
engineer
car
```

Figure 4: Program Execution with sample input

7 Improvements and Extensions

The algorithm that is given in project description has some unnecessary steps. At the first step all word mapped with 1. I think while we are sending to *slaves* from *master*, we can mapped with 1. Also before sending them back to *master*, we can sort them. Because when divide mapped ones to *slaves*, they get some words with words that come as raw string.

8 Difficulties Encountered

At the beginning, I tried to develop application with *Open MPI* library. But it's very low level library. Programming interface isn't very modern. I don't like to get output via parameter not returning variable. Also it has limited primitives types that you can send via message. Boost MPI has more modern interface. It's more useful. Learning 4-5 functions was enough for this project. Also *std::merge* has very interesting signature. I had difficulties while using it.

9 Conclusion

We are learning *Parallel Programming* paradigm in *CmpE344*, *CmpE322* and *CmpE300*. But all of them are theoretical. This project was a great chance to see real applications of it. Now, more stones fit to holes.

10 Appendices

Next lines contain source code of application.

```
1 #include <boost/mpi.hpp>
2 #include <iostream>
3 #include <vector>
4 #include <fstream>
5
6 /**
7  * It's rank of the master processor
8  */
9 #define MASTER_PROCESSOR 0
10
11 using namespace std;
12 namespace mpi = boost::mpi;
13
14 /**
15  * Word struct keeps information of string and it's count
16  */
17 class Word {
18 public:
19     string text;
20     int count;
21
22     /**
23      * Default constructor
24      */
25     Word() {
26         text = "";
27         count = 0;
28     }
29
30     /**
31      * Default constructor
32      */
33     Word(string text_, int count_) {
34         text = text_;
35         count = count_;
36     }
37
```



```

38     /**
39      * Override equal operator
40      * @param w Word to compare with this one
41      * @return true if text of the this word is equal to
42      ↪ other one
43      */
44     bool operator==(Word w) const {
45         return text == w.text;
46     }
47 private:
48     friend class boost::serialization::access;
49     /**
50      * Serializes Word struct for sending via Boost MPI
51      */
52     template<class Archive>
53     void serialize(Archive &ar, const unsigned int version) {
54         ar & text;
55         ar & count;
56     }
57 };
58
59 /**
60  * Compare two words depends on their text
61  */
62 struct wordComparator {
63     bool operator()(Word w1, Word w2) {
64         return w1.text < w2.text;
65     }
66 };
67
68 /**
69  * Entry point of my application
70  */
71 int main(int argc, char *argv[]) {
72     mpi::environment env(argc, argv);
73     mpi::communicator world;
74
75     /**
76      * Master processor work
77      */

```

```

78     if (world.rank() == MASTER_PROCESSOR) {
79         // Contains raw strings
80         vector<string> strings;
81
82         // Checks arguments count for input and output file
83         if (argc < 2) {
84             // If no input file name is give, raise an error
85             cout << "Run the code with the following command:
86                 ↪ mpirun --oversubscribe -np [processor_count]
87                 ↪ main.o [input_file] [output_file]?" << endl;
88             return 1;
89         } else {
90             // Read raw strings
91             ifstream in(argv[1]);
92             string line;
93             while (in >> line) {
94                 strings.push_back(line);
95             }
96             in.close();
97             // If output file is given, print out to it
98             if (argc == 3) {
99                 freopen(argv[2], "w", stdout);
100             }
101
102             // Divide raw string array and send to slaves
103             for (int i = 1; i < world.size(); i++) {
104                 vector<string> substrings;
105                 for (int j = 0; j < strings.size(); j++) {
106                     if (j % (world.size() - 1) == i - 1)
107                         ↪ substrings.push_back(strings[j]);
108                 }
109                 world.send(i, 1, substrings);
110             }
111
112             // Collect mapped word struct vectors from slaves
113             vector<Word> words;
114             for (int i = 1; i < world.size(); ++i) {
115                 vector<Word> subwords;
116                 world.recv(i, 2, subwords);
117                 words.insert(words.end(), subwords.begin(),
118                     ↪ subwords.end());

```

```

115     }
116     // Divide mapped word vector and send to slaves
117     for (int i = 1; i < world.size(); i++) {
118         vector<Word> subwords;
119         for (int j = 0; j < words.size(); j++) {
120             if (j % (world.size() - 1) == i - 1)
121                 ↪ subwords.push_back(words[j]);
122         }
123         world.send(i, 3, subwords);
124     }
125     // Collect sorted mapped word struct vectors from
126     ↪ slaves and merge
127     vector<Word> sorted_words;
128     for (int i = 1; i < world.size(); ++i) {
129         vector<Word> subwords;
130         world.recv(i, 4, subwords);
131         vector<Word> temp_words(sorted_words.size() +
132             ↪ subwords.size());
133         // Merge sorted vectors that come from slaved
134         merge(sorted_words.begin(), sorted_words.end(),
135             ↪ subwords.begin(), subwords.end(),
136             ↪ temp_words.begin(), wordComparator());
137         swap(sorted_words, temp_words);
138     }
139     // Reduce words vector and sum up their counts
140     vector<Word> reduced_words;
141     Word last_word = sorted_words.front();
142     for (int i = 1; i < sorted_words.size() + 1; i++) {
143         if (last_word == sorted_words[i]) {
144             // If the last one is same word, increase
145             ↪ count
146             last_word.count++;
147         } else {
148             // If it's different word, push it to reduced
149             ↪ one and change last word
150             reduced_words.push_back(last_word);
151             last_word = sorted_words[i];
152         }
153     }
154     // Print out strings and their counts
155     for (int i = 0; i < reduced_words.size(); i++) {

```

```

149         cout << reduced_words[i].text << " " <<
           ↪ reduced_words[i].count << endl;
150     }
151 } else {
152     /**
153      * Slave processors works
154      */
155     // Receive raw string vector
156     vector<string> strings;
157     world.recv(MASTER_PROCESSOR, 1, strings);
158     // Map string to Word with count
159     vector<Word> words;
160     for (int i = 0; i < strings.size(); i++) {
161         words.push_back(Word(strings[i], 1));
162     }
163     // Sent back to mapped struct to master
164     world.send(MASTER_PROCESSOR, 2, words);
165     // Receive mapped vectors to sort
166     vector<Word> words_to_sort;
167     world.recv(MASTER_PROCESSOR, 3, words_to_sort);
168     // Sort subvectors
169     sort(words_to_sort.begin(), words_to_sort.end(),
           ↪ wordComparator());
170     // Send sorted vectors to master
171     world.send(MASTER_PROCESSOR, 4, words_to_sort);
172 }
173 return 0;
174 }

```

Code 2: Source Code of Application